

Interpreting Multimodal Referring Expressions in Real Time

Miles Eldon¹ and Stefanie Tellex¹

Abstract—Robots that collaborate with humans must be able to identify objects used for shared tasks, for example tools such as a knife for assistance at cooking, or parts such as a screw on a factory floor. Existing work has addressed this problem in single modalities, such as natural language or gesture, but a gap remains in creating real-time multimodal systems that simultaneously fuse information from language and gesture in a principled mathematical framework. We define a multimodal Bayes’ filter for interpreting referring expressions to objects using language and gesture in real time. Our approach outputs an estimate of which object the person was referencing several times per second, enabling a robot to dynamically respond with backchannel feedback while a person is still communicating. We collected a new RGB-D and audio dataset of people referring to objects in a tabletop setting and demonstrate that our approach successfully integrates information from language and gesture in real time to quickly and accurately identify objects continuously.

I. INTRODUCTION

In order for humans and robots to collaborate in complex tasks, robots must be able to understand people’s references to objects in the external world. For example, a robotic cooking assistant might fetch ingredients and tools, while a robotic factory assistant could deliver a part or a hospital robot could deliver water to a bedridden patient; Figure 1 shows a robot handing a tool to a worker in a factory. To refer to objects, people use a combination of language, gesture, and body language such as eye gaze and looking. People provide these signals continuously, and a person’s reference can quickly change based on new information about the domain. Moreover, a human listener responds to these signals as they are given using *backchannels*, for example nodding their head when they understand and looking confused or interrupting to ask a question when they do not. ? refers to this continuous dance as *joint activity* and compares language use to playing a duet because of its collaborative nature, where both parties act to establish common ground and reduce uncertainty. Language and gesture co-occur and the relative timing of speech and gesture is critical for accurate understanding.

In contrast, many existing unimodal models that do not integrate information from language and gesture [??], even though people fluidly use language and gesture together. Approaches that fuse information from language and gesture [?] do not take into account that information appears to the system over a period of time. These approaches make it impossible for a robot to provide back-channel feedback, because of the length of time required to interpret




Fig. 1. Robots that collaborate with people need to understand their references to objects in the environment. For example, if a person asks for a tool using language and gesture, the robot needs to interpret the person’s reference in order to pick up the correct tool.

the communication and because of the inability to interpret partial utterances.

Taking into account time is important because it is the foundation of seeing language and conversation as a joint activity embedded in time. Responding quickly to a person’s input makes interaction more fluid and enables a robot to provide back-channel feedback based on its ability to understand: when it is confident, it can indicate that it is confident, and when it is unsure, it can indicate that. This backchannel feedback could elicit appropriate responses from the person: they will move to the next task when the robot understands, or provide more information to disambiguate when the robot is confused.

To provide a foundation for these capabilities, we propose a Bayes’ filtering approach for interpreting multimodal information from language and gesture [?]. Our framework relies on a factored observation probability that fuses information from language, hand gestures, and head gestures in real time to continuously estimate the object a person is referring to in the real world. We demonstrate our model in simulation, as well as providing quantitative results on a real-world RGB-D corpus of people referring to objects in the environment. These results demonstrate that our approach quickly and accurately fuses multimodal information in real time to continuously estimate the object a person is referencing.

¹Computer Science Department, Brown University

We describe how this approach can be used to program backchannel responses and demonstrate these responses on a real robot.

II. RELATED WORK

? proposed that conversation is a *joint activity*, a coordinated, collaborative processes akin to playing a duet or performing a waltz. The two participants must establish *common ground*. Common ground in dialog is distinct from symbol grounding [?], which is the problem of mapping from language to aspects of the external world. Common ground refers to the process of two conversational participants establishing joint understanding about the beliefs of the others. To establish common ground, people use *backchannel* feedback, such as head nods, looks of confusion, as well as explicit request for clarification such as asking a question. These mechanisms enable the participants in a conversation to engage in a feedback loop to iteratively establish common ground as the conversation progresses. Our approach for interpreting language and gesture in real time provides a foundation for producing backchannel feedback with a robot, pointing toward increased robustness as a person and robot iteratively establish common ground and actively communicate to reduce errors.

A large body of work focuses on language understanding for robots [????]. This work does not take into account the continuous nature of natural language input. ? presents a framework for interpreting open-domain references to objects but focuses on interpreting language rather than language combined with gesture. ? presented a framework for understanding language incrementally in real time dialog but did not use gesture and did not use a corpus-based evaluation. Our approach is related to ? but focuses on interpreting a person’s gestures rather than enabling a robot to generate pointing gestures.

Many existing approaches for interpreting gesture rely on fixed vocabularies of gesture, such as “stop” or “follow” [??] without a principled way for fusing information from language and gesture. Our work unifies language and gesture interpretation into a single mathematical framework, and focuses on parameterized gestures such as pointing.

? presented a multimodal framework for interpreting unscripted references to tabletop objects using language and gesture. Our approach similarly focuses on tabletop objects but uses language, gesture, and head pose, and integrates these disparate data sources continuously over time using a Bayes’ filtering framework. This approach enables the robot to continuously process new information and produce an estimate that converges over time to the correct object as new information is observed from the person.

LP and ICMI stuff

POMDP dialog (Steve Young cites gesture stuff.)

Matthias.

? created a framework enabling a robot to produce gesture. Similarly, ? described an approach for enabling a robot to generate language by inverting a semantics framework. Our long-term aim is that by combining these types of generation

approaches with real-time understanding, the robot will produce back-channel feedback that closes the loop of dialog and enable it to participate in dialog as a joint activity.

III. TECHNICAL APPROACH

Our aim is to estimate a distribution over the object that a person is referring to given language and gesture inputs. We frame the problem as a Bayes’ filter [?], where the hidden state, \mathcal{X} , is the set of m objects in the scene that the person is currently referencing. The robot observes the person’s actions and speech, \mathcal{Z} , and at each time step estimates a distribution over \mathcal{X} :

$$p(x_t|z_0 \dots z_{0:t}) \quad (1)$$

To estimate this distribution, we alternate performing a time update and a measurement update. The time update updates the belief that the user is referring to a specific subset of objects given previous information:

$$p(x_t|z_{0:t-1}) = \int p(x_t|x_{t-1}) \times p(x_{t-1}|z_{0:t-1}) dx_{t-1} \quad (2)$$

The measurement update combines the previous belief with the newest observation to update each belief state:

$$p(x_t|z_{0:t}) = \frac{p(z_t|x_t) \times p(x_t|z_{0:t-1})}{p(z_t|z_{0:t-1})} \quad (3)$$

$$\propto p(z_t|x_t) \times p(x_t|z_{0:t-1}) \quad (4)$$

Algorithm 1 shows pseudocode for our approach.

A. Prediction Model

We assume that a person is likely to continue referring to the same object, but at each timestep has a small probability, c , of transitioning to a different object:

$$p(x_t|x_{t-1}) = \begin{cases} 1 - c & \text{if } x_t = x_{t-1} \\ c & \text{otherwise} \end{cases} \quad (5)$$

In our experiments, c has a value of XXX. This assumption means that the robot’s certainty slowly decays over time, in the absence of corroborating information. Our framework integrates past language and gesture information but also quickly adapts to new, conflicting information because it assumes the person has changed objects.

B. Observation Model

We assume access to an observation model of the form:

$$p(z_t|x_t) \quad (6)$$

Observations consist of a tuple consisting of a person’s actions, $\langle l, r, h, s \rangle$ where:

- l represents the observed origin (l_o) and vector (l_v) for the left arm.
- r represents the observed origin (r_o) and vector (r_v) for the right arm .
- h represents the observed origin (h_o) and vector (h_v) for head.
- s represents the observed speech from the user, consisting of a list of words.

Formally, we have:

$$p(z_t|x_t) = p(l, r, h, s|x_t) \quad (7)$$

We factor assuming that each modality is independent of the others given the state (the true object that the person is referencing):

$$p(z_t|x_t) = p(l|x_t) \times p(r|x_t) \times p(h|x_t) \times p(s|x_t) \quad (8)$$

The following sections describe how we model each type of input from the person.

Gesture. We model pointing gestures as a vector through three dimensional space. We calculate the angle between the gesture vector and the vector from the gesture origin to the mean of each cluster, and then use the PDF of a Gaussian (\mathcal{N}) with variance (σ) to determine the weight that should be assigned to that object. We define a function $A(o, p_1, p_2)$ as the angle between the two points, p_1 and p_2 with the given origin, o . Then

$$p(l|x_t) \propto \mathcal{N}(\mu_l = 0, \sigma_l, A(l_o, l_v, x_t)) \quad (9)$$

$$p(r|x_t) \propto \mathcal{N}(\mu_r = 0, \sigma_r, A(r_o, r_v, x_t)) \quad (10)$$

If the person’s arm is more than a certain angle away from the table, we assume they are referring to none of the objects, and perform an update. As a result, these gestures do not effect the robot’s estimate of the objects being referenced.

Head Pose. Head pose is modeled in the same manner as arm gestures.

$$p(h|x_t) \propto \mathcal{N}(\mu_h = 0, \sigma_h, A(h_o, h_v, x_t)) \quad (11)$$

Speech. We model speech as a bag of words. We take the words in a given speech input and count how many words in this text match descriptors (denoted x_d) of specific objects.

$$p(s|x_t) = \prod_{w \in s} p(w|x_t) \quad (12)$$

C. Null Words and Gesture

To account for continuous gesture and non-continuous speech input, we have both null poses and speech. When no words are spoken, we assume a null word which has a uniform distribution over the objects. This effect means that spoken words cause a discrete bump in probability according to the language model, which then decays over time. While gesture remains a continuous input throughout the entire interaction, many gestures have little or no meaning. To allow for these without overloading the model with noise, we also calculate the angle between each arm vector and each foot. If the angle between the arms and a foot is smaller than the angle between the arms and any object, we assume that the user is simply in a “resting” pose, and don’t incorporate that information from that arm into that update iteration. This creates a more robust model, rather than simply defaulting to the lowest vertical object whenever a user puts their arms to their sides.

D. Model Parameters

We generated model parameters by hand, using trial and error with several subjects to fine tune them. We would gesture at an object while observing the confidence distribution and compare that to our personal belief of accuracy. We then lowering the transition probability if we felt confidence dropped too fast when a gesture ended, decreasing variance if we felt that confidence didn’t increase fast enough when we gestured. When we agreed with the control flow of the program, we committed those parameters. We considered collecting and annotating a data set to train the model parameters, however with the subjectivity of gesture, we found our initial process both accurate and expedient. We generated the language model by hand, adding to it based on results of our first user studies.

Algorithm 1: Interactive Bayes Filtering Algorithm

Input: $bel(x_{t-1}), z_t$

Output: $bel(x_t)$

```

for  $x_t$  do
     $\bar{bel}(x_t) = \prod_{x_{t-1}} p(x_t|x_{t-1}) * bel(x_{t-1})$ 
    if not is_null_gesture(l)
         $\bar{bel}(x_t) = p(l|x_t) * \bar{bel}(x_t)$ 
    if not is_null_gesture(r)
         $\bar{bel}(x_t) = p(r|x_t) * \bar{bel}(x_t)$ 
    if not is_null_gesture(h)
         $\bar{bel}(x_t) = p(h|x_t) * \bar{bel}(x_t)$ 
    for  $w \in s$  do
         $\bar{bel}(x_t) = p(w|x_t) * \bar{bel}(x_t)$ 
    end
     $bel(x_t) = \bar{bel}(x_t)$ 
end

```

Figure 2 shows an example of the system’s execution. The person’s gesture is ambiguous, but information from the language, that is itself ambiguous (because there are two bowls in the scene, enables the system to infer the correct object. Although in this example we are demonstrating the approach at two specific timesteps, the system is updating its distribution continuously, enabling it to fuse language and gesture as it occurs and quickly updating in response to new input from the person, verbal or nonverbal. Our approach runs at 14Hz, including a 30Hz sleep cycle, on an Asus machine with 8 2.4 GHz Intel Cores that is also performing all perceptual and network processing, continuously outputting an estimate of the object the person is referencing.

IV. EVALUATION

We evaluate our model in simulation, comparing the full model to versions without multimodal information. Additionally we assessed its performance on an RGB-D audio

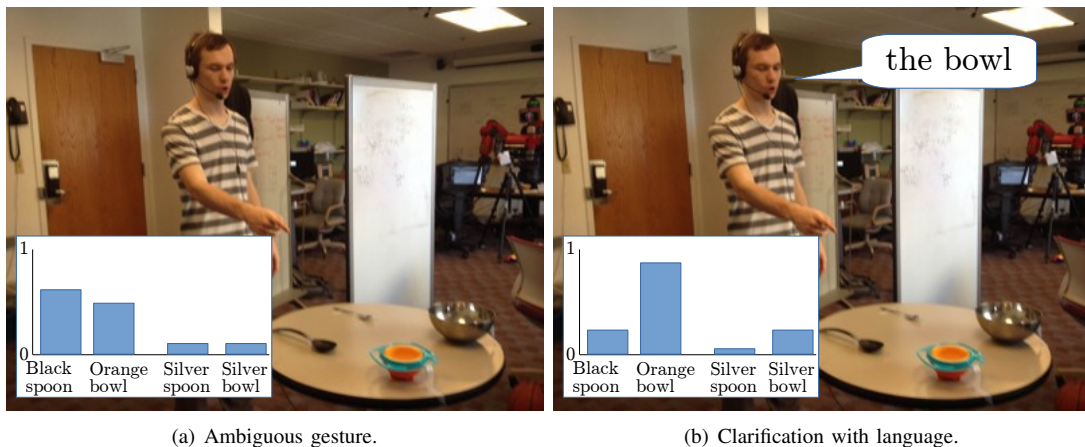


Fig. 2. After an ambiguous gesture, the model has a uniform distribution between two objects (a). Clarification with language causes a probabilistic update leaving the model highly confident it has inferred the correct object (b).

TABLE I
SIMULATION RESULTS

	$\sigma^2 = 0.5$	$\sigma^2 = 1.0$
Language alone	36.5%	37%
Head alone	50.9%	35.8%
Arms alone	62.4%	42.1%
Language and gesture	65.7%	54.1%

and video corpus of people referring to objects. Finally we created an end-to-end robotic demonstration, demonstrated in the video attachment to our paper, and available online¹.

A. Simulation Results

We evaluated our approach in simulation by generating data from the model and assessing its accuracy at estimating the object being referenced. We generate simulated pointing data for each hand, the face, as well as spoken utterances at each timestep according to the model parameters. We then use these parameters to update the system’s estimate of the object being referred to. Table I shows the results. Our accuracy metric is the fraction of time that the robot is pointing at the correct object. We report performance using language alone, gesture alone, and language combined with gesture. These results demonstrate that the system is successfully able to fuse multi-modal information to achieve higher accuracy than each modality alone, but do not show performance in the real world. While a gesture and language output probabilistically referencing the true object are produced at each time step in simulation, the multimodal accuracy is robust to severe increases in single-modal noise, as you can see in the table when we increase the sample variance.

B. Real-World Corpus-Based Results

Our real-world experiments measured our algorithm’s performance when a person referred to an object visually and with gesture. The subject stood in front of a table with four objects placed approximately one foot apart, forming four

Fig. 3. Scene from our dataset.

TABLE II
REAL-WORLD RESULTS

Random	25%
Language only	32.7% +/- 3.5%
Gesture only	73.31% +/- 4%
Head only	23.4% +/- 2%
Multimodal (Language and Gesture)	80.52% +/- 2%
Multimodal (All)	64.3% +/- 3%

corners of a square, as shown in Figure 3. We instructed subjects to ask for the indicated object in the most naturally way possible, using whatever combination of gesture and language they felt was appropriate. We indicted the object to refer to using a laser pointer, and we periodically shifted to a different object on a predetermined schedule. They wore a microphone to pick up high-quality audio, and we tracked their body pose using the NITE tracker [?]. We used a single Kinect that was able to capture both the user and the table.

We used the HTML5 Speech Recognition package in conjunction with Google Chrome to recognize speech. This package reports incremental output in real time as recognition proceeds, and we perform a model update each time a new word is perceived. Our training procedure works on actual speech recognition results rather than transcript speech, enabling the algorithm to adapt to the errors produced by the recognizer: if it returns “cake” instead of “cup,” our language model will correctly associate the word “cake” with the “cup” object.

Results showing the percent of the time the estimated most likely object was the true object appear in Table II with 95% confidence intervals. Additionally, we display accuracy at the end of an interaction in III.

¹video reference

TABLE III
REAL-WORLD RESULTS (END OF INTERACTION)

Random	25%
Language only	46.15%
Gesture only	80.0%
Head only	18.46%
Multimodal (Language and Gesture)	90.77%
Multimodal (All)	61.54%

C. Robotic Demonstration

Because our approach enables a robot to quickly and constantly monitor a person’s references to an object, a robot can respond to these estimates in real time. We demonstrate this behavior by enabling Baxter to demonstrate its certainty about what object is being referenced, this eliciting more feedback from the person. When the robot is very unsure, its arm moves back and forth between the candidate objects. When it is sure, then it moves with more precision. The video shows this behavior, as a person provides more information about what object is being referred to by the person.

V. CONCLUSION

We have demonstrated a Bayes’ filtering approach to interpreting a person’s multimodal language and gesture references to objects continuously in real time. Our approach enables a robot to understand a person’s references to objects in the real world.

In the future we plan to expand our language model to incorporate models of compositional semantics and lower-level visual features so that the robot is not limited to pre-specified object models. Additionally we aim to enable the robot to generate back-channel feedback based on its model. ? created a framework for generating legible gesture, and we anticipate that enabling a robot to respond by pointing as in ? when it is sure and reflecting its confusing when it is unsure. Closing the loop will enable the human-robot dyad to increase efficiency and enable the robot to accurately infer the human’s intent, naturally eliciting more information when it is confused and indicating that it has understood when it is sure. This paper represents steps toward continuous language understanding and the vision presented by ? of language as joint activity.